

```
In [*]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

Load the dataset

```
In [*]: df=pd.read_csv('abalone.csv')
```

```
In [3]: df
```


Out[3]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.1500	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.0700	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.2100	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.1550	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.0550	7
...	...	...	...	...	...	...	...	...	...
4172	F	0.565	0.450	0.165	0.8870	0.3700	0.2390	0.2490	11
4173	M	0.590	0.440	0.135	0.9660	0.4390	0.2145	0.2605	10
4174	M	0.600	0.475	0.205	1.1760	0.5255	0.2875	0.3080	9
4175	F	0.625	0.485	0.150	1.0945	0.5310	0.2610	0.2960	10
4176	M	0.710	0.555	0.195	1.9485	0.9455	0.3765	0.4950	12


4177 rows × 9 columns

Univariate analysis

```
In [4]: import seaborn as sns  
plt.scatter(df.index, df['Sex'])  
plt.show()
```




```
In [5]: plt.scatter(df.index, df['Length'])  
plt.show()
```




```
In [6]: sns.stripplot(y=df['Diameter'])
```

```
Out[6]: <AxesSubplot:ylabel='Diameter'>
```



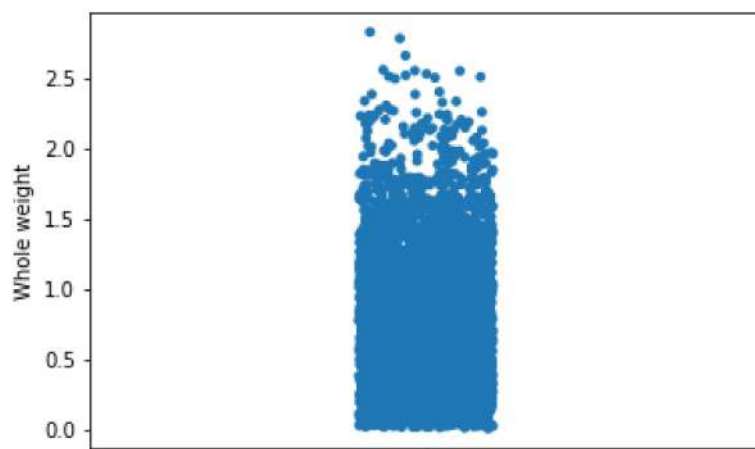
```
In [7]: sns.stripplot(y=df['Height'])
```

```
Out[7]: <AxesSubplot:ylabel='Height'>
```




```
In [8]: sns.stripplot(y=df['Whole weight'])
```

```
Out[8]: <AxesSubplot:ylabel='Whole weight'>
```




```
In [9]: plt.scatter(df.index,df['Shucked weight'])  
plt.show()
```




```
In [10]: sns.stripplot(y=df['Viscera weight'])
```

```
Out[10]: <AxesSubplot:ylabel='Viscera weight'>
```




```
In [11]: sns.stripplot(y=df['Shell weight'])
```

```
Out[11]: <AxesSubplot:ylabel='Shell weight'>
```




```
In [12]: plt.scatter(df.index,df['Rings'])  
plt.show()
```




Bi-variate analysis

```
In [13]: sns.boxplot(x='Sex',y='Rings',data=df)  
plt.show()
```




```
In [14]: sns.violinplot(x="Height",y="Rings",size=8,data=df)  
plt.show()
```




```
In [15]: plt.plot(df["Length"],df["Rings"],'o')  
plt.ylabel("Rings")  
plt.xlabel("Length")
```

Out[15]: Text(0.5, 0, 'Length')




```
In [16]: plt.plot(df["Diameter"],df["Rings"],'o')  
plt.ylabel("Rings")  
plt.xlabel("Diameter")
```

Out[16]: Text(0.5, 0, 'Diameter')




```
In [17]: plt.plot(df["Whole weight"],df["Rings"],'.')  
plt.ylabel("Rings")  
plt.xlabel("Length")
```

Out[17]: Text(0.5, 0, 'Length')



```
In [18]: plt.plot(df["Shell weight"],df["Rings"],'*')  
plt.ylabel("Rings")  
plt.xlabel("Length")
```


Out[18]: Text(0.5, 0, 'Length')






```
In [19]: plt.plot(df["Shucked weight"],df["Rings"],'o')  
plt.ylabel("Rings")  
plt.xlabel("Length")
```

```
Out[19]: Text(0.5, 0, 'Length')
```




```
In [20]: plt.plot(df["Viscera weight"],df["Rings"],'.')  
plt.ylabel("Rings")  
plt.xlabel("Length")
```

```
Out[20]: Text(0.5, 0, 'Length')
```



Multivariate Analysis

```
In [21]: sns.pairplot(df,hue="Rings",height=3)  
plt.show()
```



Descriptive analysis

```
In [22]: df.describe().T
```

Out[22]:

	count	mean	std	min	25%	50%	75%	max
Length	4177.0	0.523992	0.120093	0.0750	0.4500	0.5450	0.615	0.8150
Diameter	4177.0	0.407881	0.099240	0.0550	0.3500	0.4250	0.480	0.6500
Height	4177.0	0.139516	0.041827	0.0000	0.1150	0.1400	0.165	1.1300
Whole weight	4177.0	0.828742	0.490389	0.0020	0.4415	0.7995	1.153	2.8255
Shucked weight	4177.0	0.359367	0.221963	0.0010	0.1860	0.3360	0.502	1.4880
Viscera weight	4177.0	0.180594	0.109614	0.0005	0.0935	0.1710	0.253	0.7600
Shell weight	4177.0	0.238831	0.139203	0.0015	0.1300	0.2340	0.329	1.0050
Rings	4177.0	9.933684	3.224169	1.0000	8.0000	9.0000	11.000	29.0000

Handling with missing data

```
In [23]: df = pd.DataFrame(df)
df.isnull()
```

Out[23]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False
...	...	...	...	...	...	...	...	...	...
4172	False	False	False	False	False	False	False	False	False
4173	False	False	False	False	False	False	False	False	False
4174	False	False	False	False	False	False	False	False	False
4175	False	False	False	False	False	False	False	False	False
4176	False	False	False	False	False	False	False	False	False

4177 rows × 9 columns

In [24]: `df.fillna(0)`

Out[24]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.1500	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.0700	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.2100	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.1550	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.0550	7
...	...	...	...	...	...	...	...	...	...
4172	F	0.565	0.450	0.165	0.8870	0.3700	0.2390	0.2490	11
4173	M	0.590	0.440	0.135	0.9660	0.4390	0.2145	0.2605	10
4174	M	0.600	0.475	0.205	1.1760	0.5255	0.2875	0.3080	9
4175	F	0.625	0.485	0.150	1.0945	0.5310	0.2610	0.2960	10
4176	M	0.710	0.555	0.195	1.9485	0.9455	0.3765	0.4950	12

4177 rows × 9 columns

In [25]: `df.isnull().sum()`


Out[25]: Sex 0  
Length 0  
Diameter 0  
Height 0  
Whole weight 0  
Shucked weight 0  
Viscera weight 0  
Shell weight 0  
Rings 0  
dtype: int64

Outliers in each attribute

```
In [26]: sns.boxplot(df['Length'],data=df)
```

C:\Users\ABI\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(


```
Out[26]: <AxesSubplot:xlabel='Length'>
```



```
In [27]: sns.boxplot(df['Diameter'],data=df)
```

C:\Users\ABI\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(


```
Out[27]: <AxesSubplot:xlabel='Diameter'>
```



```
In [28]: sns.boxplot(df['Height'],data=df)
```

C:\Users\ABI\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(

```
Out[28]: <AxesSubplot:xlabel='Height'>
```




```
In [29]: sns.boxplot(df['Whole weight'],data=df)
```

C:\Users\ABI\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

```
Out[29]: <AxesSubplot:xlabel='Whole weight'>
```






```
In [30]: sns.boxplot(df['Shucked weight'],data=df)
```

C:\Users\ABI\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```


```
Out[30]: <AxesSubplot:xlabel='Shucked weight'>
```



```
In [31]: sns.boxplot(df['Viscera weight'],data=df)
```

C:\Users\ABI\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(


```
Out[31]: <AxesSubplot:xlabel='Viscera weight'>
```



```
In [32]: sns.boxplot(df['Shell weight'],data=df)
```

C:\Users\ABI\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(


```
Out[32]: <AxesSubplot:xlabel='Shell weight'>
```



```
In [33]: sns.boxplot(df['Rings'],data=df)
```

C:\Users\ABI\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(

```
Out[33]: <AxesSubplot:xlabel='Rings'>
```



```
In [38]: Q1 = df.quantile(0.25)
Q3 = df.quantile(0.75)
IQR = Q3-Q1
print(IQR)
```

```
Length          0.1650
Diameter        0.1300
Height          0.0500
Whole weight    0.7115
Shucked weight  0.3160
Viscera weight  0.1595
Shell weight    0.1990
Rings           3.0000
dtype: float64
```

```
In [39]: df.shape
```

```
Out[39]: (4177, 9)
```

```
In [43]: upper = np.where(df >= (Q3+1.5*IQR))
lower = np.where(df <= (Q1-1.5*IQR))
df.drop(upper[0], inplace = True)
df.drop(lower[0], inplace = True)
df.shape
```

C:\Users\ABI\AppData\Local\Temp\ipykernel\_15784\1776641308.py:1: FutureWarning: Automatic reindexing on DataFrame vs Series comparisons is deprecated and will raise ValueError in a future version. Do `left, right = left.align(right, axis=1, copy=False)` before e.g. `left == right`

```
upper = np.where(df >= (Q3+1.5*IQR))
```

C:\Users\ABI\AppData\Local\Temp\ipykernel\_15784\1776641308.py:2: FutureWarning: Automatic reindexing on DataFrame vs Series comparisons is deprecated and will raise ValueError in a future version. Do `left, right = left.align(right, axis=1, copy=False)` before e.g. `left == right`

```
lower = np.where(df <= (Q1-1.5*IQR))
```


Out[43]: (3781, 9)

```
In [44]: sns.boxplot(df['Length'],data=df)
```

C:\Users\ABI\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[44]: <AxesSubplot:xlabel='Length'>




```
In [45]: sns.boxplot(df['Rings'],data=df)
```

C:\Users\ABI\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```


```
Out[45]: <AxesSubplot:xlabel='Rings'>
```



```
In [46]: sns.boxplot(df['Height'],data=df)
```

C:\Users\ABI\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(


```
Out[46]: <AxesSubplot:xlabel='Height'>
```



```
In [47]: sns.boxplot(df['Diameter'],data=df)
```

C:\Users\ABI\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(

```
Out[47]: <AxesSubplot:xlabel='Diameter'>
```






```
In [48]: sns.boxplot(df['Whole weight'],data=df)
```

C:\Users\ABI\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```


```
Out[48]: <AxesSubplot:xlabel='Whole weight'>
```



```
In [49]: sns.boxplot(df['Shucked weight'],data=df)
```

C:\Users\ABI\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(


```
Out[49]: <AxesSubplot:xlabel='Shucked weight'>
```



```
In [50]: sns.boxplot(df['Shell weight'],data=df)
```

C:\Users\ABI\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(


```
Out[50]: <AxesSubplot:xlabel='Shell weight'>
```



```
In [51]: sns.boxplot(df['Viscera weight'],data=df)
```

C:\Users\ABI\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
warnings.warn(

```
Out[51]: <AxesSubplot:xlabel='Viscera weight'>
```



Label encoding for categorical data

```
In [53]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler

from sklearn.linear_model import LinearRegression
from sklearn.svm import SVR
from sklearn.tree import DecisionTreeRegressor
from sklearn import metrics
```

```
In [55]: le=LabelEncoder()
df['Sex']=le.fit_transform(df['Sex'])
```

In [56]: df

Out[56]:

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	2	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.1500	15
1	2	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.0700	7
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.2100	9
3	2	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.1550	10
4	1	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.0550	7
...	...	...	...	...	...	...	...	...	...
4172	0	0.565	0.450	0.165	0.8870	0.3700	0.2390	0.2490	11
4173	2	0.590	0.440	0.135	0.9660	0.4390	0.2145	0.2605	10
4174	2	0.600	0.475	0.205	1.1760	0.5255	0.2875	0.3080	9
4175	0	0.625	0.485	0.150	1.0945	0.5310	0.2610	0.2960	10
4176	2	0.710	0.555	0.195	1.9485	0.9455	0.3765	0.4950	12

3781 rows × 9 columns

Splitting the Data into dependent and Independent Variables

In [57]: `X = df.iloc[:, :-1].values`  
`y = df.iloc[:, -1].values`

Scaling independent variables

In [58]: `scaler = StandardScaler()`  
`scaler.fit(df)`

Out[58]: `StandardScaler()`

Splitting training and test data

In [59]: `train_X, val_X, train_y, val_y = train_test_split(X, y, test_size = 0.2, random_state=42)`

In [60]: `print("Shape of Training X :", train_X.shape)`  
`print("Shape of Validation X :", val_X.shape)`

Shape of Training X : (3024, 8)  
 Shape of Validation X : (757, 8)

```
In [61]: print("Shape of Training y :",train_y.shape)
print("Shape of Validation y :",val_y.shape)
```

```
Shape of Training y : (3024,)
Shape of Validation y : (757,)
```

Linear regression

```
In [62]: lr = LinearRegression()
lr.fit(train_X,train_y)
```

```
Out[62]: LinearRegression()
```

```
In [63]: %%time
y_pred_val_lr = lr.predict(val_X)
print('MAE on Validation set :',metrics.mean_absolute_error(val_y, y_pred_val_lr))
print("\n")
print('MSE on Validation set :',metrics.mean_squared_error(val_y, y_pred_val_lr))
print("\n")
print('RMSE on Validation set :',np.sqrt(metrics.mean_squared_error(val_y, y_pred_val_lr)))
print("\n")
print('R2 Score on Validation set :',metrics.r2_score(val_y, y_pred_val_lr))
print("\n")
```

```
MAE on Validation set : 1.2719689486359296
```

```
MSE on Validation set : 2.7606215450501024
```

```
RMSE on Validation set : 1.1278160083257949
```

```
R2 Score on Validation set : 0.5119499107890585
```

```
Wall time: 4.82 ms
```

Support vector machine

```
In [64]: svm = SVR()
svm.fit(train_X,train_y)
```

```
Out[64]: SVR()
```

```
In [65]: %%time
y_pred_val_svm = svm.predict(val_X)
print('MAE on Validation set :',metrics.mean_absolute_error(val_y, y_pred_val_svm))
print("\n")
print('MSE on Validation set :',metrics.mean_squared_error(val_y, y_pred_val_svm))
print("\n")
print('RMSE on Validation set :',np.sqrt(metrics.mean_absolute_error(val_y, y_pred_val_svm)))
print("\n")
print('R2 Score on Validation set :',metrics.r2_score(val_y, y_pred_val_svm))
print("\n")
```

MAE on Validation set : 1.220895278727089

MSE on Validation set : 2.7012620714060267

RMSE on Validation set : 1.104941301032362

R2 Score on Validation set : 0.5224440679687887

Wall time: 496 ms

Decision tree regressor

```
In [66]: dc = DecisionTreeRegressor(random_state = 0)
dc.fit(train_X,train_y)
```

Out[66]: DecisionTreeRegressor(random\_state=0)

```
In [67]: %%time
y_pred_val_dc = dc.predict(val_X)
print('MAE on Validation set :',metrics.mean_absolute_error(val_y, y_pred_val_dc))
print("\n")
print('MSE on Validation set :',metrics.mean_squared_error(val_y, y_pred_val_dc))
print("\n")
print('RMSE on Validation set :',np.sqrt(metrics.mean_absolute_error(val_y, y_pred_val_dc)))
print("\n")
print('R2 Score on Validation set :',metrics.r2_score(val_y, y_pred_val_dc))
print("\n")
```

MAE on Validation set : 1.6393659180977542

MSE on Validation set : 4.88110964332893

RMSE on Validation set : 1.2803772561623212

R2 Score on Validation set : 0.13706896870869845

Wall time: 6.78 ms

```
In [ ]: Overview of R2 scores of all models
```

```
In [68]: print('Logistic Regression R2 Score on Validation set :',metrics.r2_score(val_y,
print('SVR R2 Score on Validation set :',metrics.r2_score(val_y, y_pred_val_svm))
print('Decision Tree Regressor R2 Score on Validation set :',metrics.r2_score(val_y, y_pred_val_dt))
```

Logistic Regression R2 Score on Validation set : 0.5119499107890585

SVR R2 Score on Validation set : 0.5224440679687887

Decision Tree Regressor R2 Score on Validation set : 0.13706896870869845

```
In [ ]:
```